

Lab 2 - Exercise 2

1 Experimental Setup

The following C program is used in this experiment:

```
for (int i = 0; i < N; i++) {  
    x = a * b + x;  
    y = a * b + y;  
}
```

The loop performs two independent floating-point accumulation streams. The program is compiled and executed using:

- `-O0`: no compiler optimizations
- `-O2`: aggressive compiler optimizations

Execution time is measured using `clock()`.

2 Question 1: Execution Time Comparison

The program was compiled and executed as follows:

```
gcc -O0 exercice3.c -o ex_O0  
gcc -O2 exercice3.c -o ex_O2
```

The measured execution times are shown in Table 1.

Version	Optimization Level	Time (s)
Original code	-O0	0.297
Original code	-O2	0.086

Table 1: Measured execution times

The speedup achieved by `-O2` over `-O0` is:

$$\text{Speedup} = \frac{0.297}{0.086} \approx 3.45$$

Answer: Compilation with `-O2` reduces execution time by approximately **3.45×** compared to `-O0`.

3 Question 2: Optimizations Performed at `-O2`

Analysis of the generated assembly code reveals several important optimizations at `-O2`:

3.1 Redundant Computation Elimination

The expression `a*b` is constant inside the loop. At `-O0`, it is recomputed for each update of `x` and `y`. At `-O2`, the compiler computes it once and reuses the result.

3.2 Register Allocation

At `-O0`, variables such as `x`, `y`, and the loop counter are frequently loaded from and stored to memory. At `-O2`, these variables are kept in registers, significantly reducing memory access latency.

3.3 Loop Transformations

The optimized assembly shows that the compiler performs loop unrolling. Each loop iteration processes multiple updates, reducing:

- the number of branch instructions,
- loop counter updates,
- condition checks.

3.4 Instruction Scheduling and ILP

Independent floating-point additions are reordered and scheduled to overlap execution. This allows the CPU pipeline to remain busy and hides instruction latency, increasing instruction-level parallelism.

Answer: Compared to `-O0`, `-O2` applies redundant computation elimination, register allocation, loop unrolling, and instruction scheduling to exploit ILP and reduce execution overhead.

4 Question 3: Manual Optimization vs Compiler Optimization

A manually optimized version of the code was implemented by computing `a*b` once before the loop:

```
const double t = a * b;
for (int i = 0; i < N; i++) {
    x += t;
    y += t;
}
```

This version was compiled with `-O0` and executed multiple times.

4.1 Measured Results

Average execution time for the manually optimized version:

$$T_{\text{manual}, -O0} \approx 0.292 \text{ s}$$

Version	Optimization Level	Time (s)
Original code	-O0	0.297
Manual optimized code	-O0	0.292
Original code	-O2	0.086

Table 2: Comparison of manual and automatic optimizations

4.2 Analysis

Manual optimization yields a small improvement of about 2% compared to the original `-O0` version. However, the `-O2` optimized code remains approximately **3.4× faster**.

This demonstrates that while manual optimizations can remove obvious inefficiencies, the compiler applies multiple advanced optimizations simultaneously that are difficult to reproduce manually at the source-code level.

Answer: Manual optimization at `-O0` provides only marginal improvement, whereas compiler optimization at `-O2` achieves substantially better performance.

5 Conclusion

This experiment demonstrates the effectiveness of compiler optimizations in exploiting instruction-level parallelism. While manual code restructuring can slightly improve performance, modern compilers at `-O2` outperform manual optimizations by combining register allocation, loop transformations, and instruction scheduling.